

In [2]:

```
!pip install pmdarima
```

```
Requirement already satisfied: pmdarima in c:\users\kurno\anaconda3\lib\site-packages (1.8.5)
Requirement already satisfied: joblib>=0.11 in c:\users\kurno\anaconda3\lib\site-packages (from pmdarima) (0.16.0)
Requirement already satisfied: urllib3 in c:\users\kurno\anaconda3\lib\site-packages (from pmdarima) (1.25.9)
Requirement already satisfied: scipy>=1.3.2 in c:\users\kurno\anaconda3\lib\site-packages (from pmdarima) (1.5.0)
Requirement already satisfied: numpy>=1.19.3 in c:\users\kurno\anaconda3\lib\site-packages (from pmdarima) (1.23.2)
Requirement already satisfied: setuptools!=50.0.0,>=38.6.0 in c:\users\kurno\anaconda3\lib\site-packages (from pmdarima) (49.2.0.post20200714)
Requirement already satisfied: pandas>=0.19 in c:\users\kurno\anaconda3\lib\site-packages (from pmdarima) (1.0.5)
Requirement already satisfied: Cython!=0.29.18,>=0.29 in c:\users\kurno\anaconda3\lib\site-packages (from pmdarima) (0.29.21)
Requirement already satisfied: scikit-learn>=0.22 in c:\users\kurno\anaconda3\lib\site-packages (from pmdarima) (0.23.1)
Requirement already satisfied: statsmodels!=0.12.0,>=0.11 in c:\users\kurno\anaconda3\lib\site-packages (from pmdarima) (0.11.1)
Requirement already satisfied: pytz>=2017.2 in c:\users\kurno\anaconda3\lib\site-packages (from pandas>=0.19->pmdarima) (2020.1)
Requirement already satisfied: python-dateutil>=2.6.1 in c:\users\kurno\anaconda3\lib\site-packages (from pandas>=0.19->pmdarima) (2.8.1)
Requirement already satisfied: threadpoolctl>=2.0.0 in c:\users\kurno\anaconda3\lib\site-packages (from scikit-learn>=0.22->pmdarima) (2.1.0)
Requirement already satisfied: patsy>=0.5 in c:\users\kurno\anaconda3\lib\site-packages (from statsmodels!=0.12.0,>=0.11->pmdarima) (0.5.1)
Requirement already satisfied: six>=1.5 in c:\users\kurno\anaconda3\lib\site-packages (from python-dateutil>=2.6.1->pandas>=0.19->pmdarima) (1.15.0)
```

In [3]:

```
import os
import warnings
warnings.filterwarnings('ignore')
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
plt.style.use('fivethirtyeight')
from pylab import rcParams
rcParams['figure.figsize'] = 10, 6
from statsmodels.tsa.stattools import adfuller
from statsmodels.tsa.seasonal import seasonal_decompose
from statsmodels.tsa.arima_model import ARIMA
from pmdarima.arima import auto_arima

from sklearn.metrics import mean_squared_error, mean_absolute_error
import math

for dirname, _, filenames in os.walk('https://www.kaggle.com/code/nageshsingh/stock-market-forecasting-arima/data'):
    for filename in filenames:
        print(os.path.join(dirname, filename))
```

In [5]:

```

dateparse = lambda dates: pd.datetime.strptime(dates, '%Y-%m-%d')
stock_data = pd.read_csv('Downloads/acgl.us.txt', sep=',', index_col='Date', parse_dates=
=['Date'], date_parser=dateparse).fillna(0)
stock_data

```


Out[5]:

	Open	High	Low	Close	Volume	OpenInt
Date						
2005-02-25	13.583	13.693	13.430	13.693	156240	0
2005-02-28	13.697	13.827	13.540	13.827	370509	0
2005-03-01	13.780	13.913	13.720	13.760	224484	0
2005-03-02	13.717	13.823	13.667	13.810	286431	0
2005-03-03	13.783	13.783	13.587	13.630	193824	0
...
2017-11-06	94.490	95.650	94.020	95.550	420192	0
2017-11-07	95.860	95.950	95.200	95.560	464011	0
2017-11-08	95.410	95.900	94.890	95.450	471756	0
2017-11-09	94.930	96.140	94.470	95.910	353498	0
2017-11-10	95.890	95.990	94.390	95.350	452833	0

3201 rows × 6 columns

In [11]:

```
plt.figure(figsize=(10,6))  
plt.grid(True)  
plt.xlabel('Date')  
plt.ylabel('Close Prices')  
plt.plot(stock_data['Close'])  
plt.title('ARCH CAPITAL GROUP closing price')  
plt.show()
```



In [17]:


```
df_close=stock_data['Close']
```

In [18]:

```
df_close.plot(kind='kde')
```


Out[18]:

<matplotlib.axes._subplots.AxesSubplot at 0x1bd775f3c70>



In [19]:

```
def test_stationarity(timeseries):  
    #Determining rolling statistics  
    rolmean = timeseries.rolling(12).mean()  
    rolstd = timeseries.rolling(12).std()  
    #Plot rolling statistics:  
    plt.plot(timeseries, color='blue',label='Original')  
    plt.plot(rolmean, color='red', label='Rolling Mean')  
    plt.plot(rolstd, color='black', label = 'Rolling Std')  
    plt.legend(loc='best')  
    plt.title('Rolling Mean and Standard Deviation')  
    plt.show(block=False)  
  
    print("Results of dickey fuller test")  
    adft = adfuller(timeseries,autolag='AIC')  
    # output for dft will give us without defining what the values are.  
    #hence we manually write what values does it explains using a for loop  
    output = pd.Series(adft[0:4],index=['Test Statistics','p-value','No. of lags used',  
    'Number of observations used'])  
    for key,values in adft[4].items():  
        output['critical value (%s)'%key] = values  
    print(output)  
  
test_stationarity(df_close)
```



Results of dickey fuller test


Test Statistics	1.374899
p-value	0.996997
No. of lags used	5.000000
Number of observations used	3195.000000
critical value (1%)	-3.432398
critical value (5%)	-2.862445
critical value (10%)	-2.567252

dtype: float64

In [20]:

```
result = seasonal_decompose(df_close, model='multiplicative', freq = 30)
fig = plt.figure()
fig = result.plot()
fig.set_size_inches(16, 9)
```


<Figure size 720x432 with 0 Axes>



In [21]:

```
from pylab import rcParams
rcParams['figure.figsize'] = 10, 6
df_log = np.log(df_close)
moving_avg = df_log.rolling(12).mean()
std_dev = df_log.rolling(12).std()
plt.legend(loc='best')
plt.title('Moving Average')
plt.plot(std_dev, color="black", label = "Standard Deviation")
plt.plot(moving_avg, color="red", label = "Mean")
plt.legend()
plt.show()
```

No handles with labels found to put in legend.




In [22]:

```
train_data, test_data = df_log[3:int(len(df_log)*0.9)], df_log[int(len(df_log)*0.9):]  
plt.figure(figsize=(10,6))  
plt.grid(True)  
plt.xlabel('Dates')  
plt.ylabel('Closing Prices')  
plt.plot(df_log, 'green', label='Train data')  
plt.plot(test_data, 'blue', label='Test data')  
plt.legend()
```

Out[22]:

<matplotlib.legend.Legend at 0x1bd7be0d670>



In [23]:

```
model_autoARIMA = auto_arima(train_data, start_p=0, start_q=0,
                             test='adf',          # use adftest to find optimal 'd'
                             max_p=3, max_q=3,    # maximum p and q
                             m=1,                # frequency of series
                             d=None,              # let model determine 'd'
                             seasonal=False,      # No Seasonality
                             start_P=0,
                             D=0,
                             trace=True,
                             error_action='ignore',
                             suppress_warnings=True,
                             stepwise=True)
print(model_autoARIMA.summary())
model_autoARIMA.plot_diagnostics(figsize=(15,8))
plt.show()
```

Performing stepwise search to minimize aic

```

ARIMA(0,1,0)(0,0,0)[0] intercept : AIC=-16491.508, Time=0.77 sec
ARIMA(1,1,0)(0,0,0)[0] intercept : AIC=-16525.993, Time=0.46 sec
ARIMA(0,1,1)(0,0,0)[0] intercept : AIC=-16527.964, Time=0.70 sec
ARIMA(0,1,0)(0,0,0)[0] : AIC=-16488.323, Time=0.18 sec
ARIMA(1,1,1)(0,0,0)[0] intercept : AIC=-16527.157, Time=1.74 sec
ARIMA(0,1,2)(0,0,0)[0] intercept : AIC=-16527.120, Time=1.75 sec
ARIMA(1,1,2)(0,0,0)[0] intercept : AIC=-16528.152, Time=3.70 sec
ARIMA(2,1,2)(0,0,0)[0] intercept : AIC=inf, Time=4.12 sec
ARIMA(1,1,3)(0,0,0)[0] intercept : AIC=-16526.183, Time=4.94 sec
ARIMA(0,1,3)(0,0,0)[0] intercept : AIC=-16524.974, Time=2.05 sec
ARIMA(2,1,1)(0,0,0)[0] intercept : AIC=-16525.435, Time=1.41 sec
ARIMA(2,1,3)(0,0,0)[0] intercept : AIC=-16516.417, Time=0.88 sec
ARIMA(1,1,2)(0,0,0)[0] : AIC=-16527.597, Time=0.95 sec

```

Best model: ARIMA(1,1,2)(0,0,0)[0] intercept

Total fit time: 23.685 seconds

SARIMAX Results

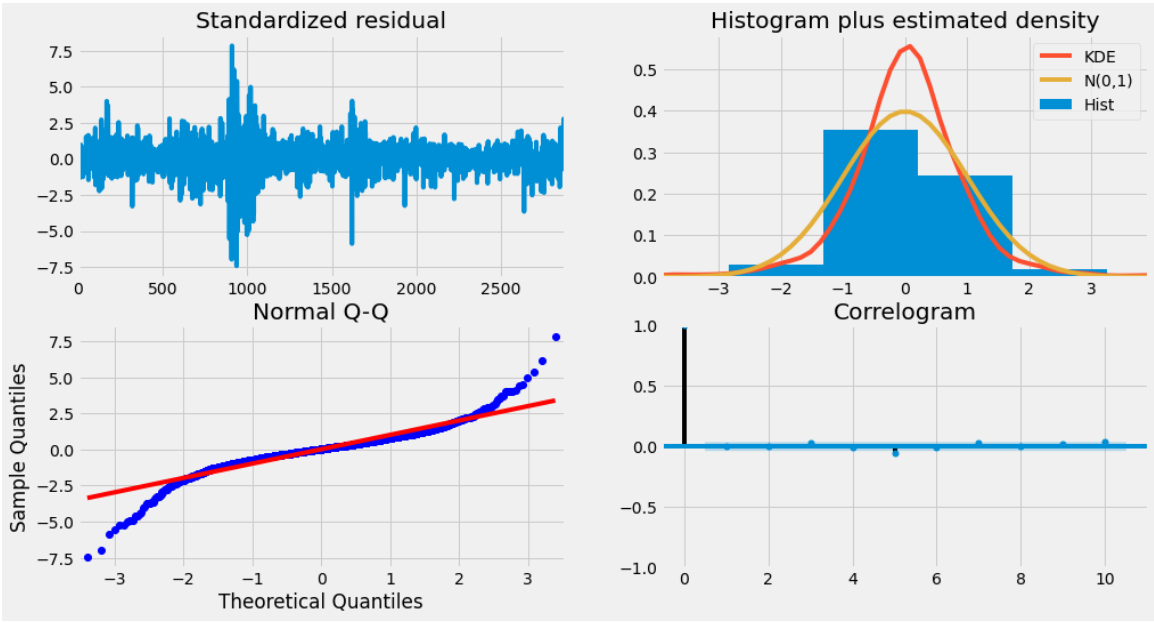
```

=====
====
Dep. Variable:          y    No. Observations:
2877
Model:                SARIMAX(1, 1, 2)    Log Likelihood          826
9.076
Date:                Thu, 18 Aug 2022    AIC                      -1652
8.152
Time:                23:10:51    BIC                      -1649
8.331
Sample:                0    HQIC                      -1651
7.402
- 2877
Covariance Type:      opg
=====
====
               coef    std err          z      P>|z|      [0.025      0.
975]
-----
----
intercept    2.27e-05    6.69e-06     3.393     0.001    9.59e-06    3.58
e-05
ar.L1         0.9538         0.009    104.122     0.000         0.936
0.972
ma.L1        -1.0708         0.015    -73.543     0.000        -1.099    -
1.042
ma.L2         0.0877         0.012     7.501     0.000         0.065
0.111
sigma2        0.0002    2.32e-06    80.812     0.000         0.000
0.000
=====
=====
Ljung-Box (Q):                121.69    Jarque-Bera (JB):
7206.78
Prob(Q):                      0.00    Prob(JB):
0.00
Heteroskedasticity (H):        0.30    Skew:
-0.39
Prob(H) (two-sided):           0.00    Kurtosis:
10.72
=====
=====

```

Warnings:

```
[1] Covariance matrix calculated using the outer product of gradients (complex-step).
```



In [24]:

```
model = ARIMA(train_data, order=(1,1,2))  
fitted = model.fit(dispatch=-1)  
print(fitted.summary())
```

```
C:\Users\kurno\anaconda3\lib\site-packages\statsmodels\tsa\base\tsa_model.py:216: ValueWarning: A date index has been provided, but it has no associated frequency information and so will be ignored when e.g. forecasting.
  warnings.warn('A date index has been provided, but it has no'
C:\Users\kurno\anaconda3\lib\site-packages\statsmodels\tsa\base\tsa_model.py:216: ValueWarning: A date index has been provided, but it has no associated frequency information and so will be ignored when e.g. forecasting.
  warnings.warn('A date index has been provided, but it has no'
```

ARIMA Model Results					
=====					
=====					
Dep. Variable:	D.Close	No. Observations:			
2876					
Model:	ARIMA(1, 1, 2)	Log Likelihood		827	
4.158					
Method:	css-mle	S.D. of innovations			
0.014					
Date:	Thu, 18 Aug 2022	AIC		-1653	
8.316					
Time:	23:11:53	BIC		-1650	
8.496					
Sample:	1	HQIC		-1652	
7.567					
=====					
=====					
	coef	std err	z	P> z	[0.025
0.975]					

const	0.0006	0.000	3.935	0.000	0.000
0.001					
ar.L1.D.Close	0.9145	0.040	22.740	0.000	0.836
0.993					
ma.L1.D.Close	-1.0351	0.045	-23.127	0.000	-1.123
-0.947					
ma.L2.D.Close	0.0848	0.022	3.820	0.000	0.041
0.128					
Roots					
=====					
=====					
	Real	Imaginary	Modulus		Freque
ncy					

AR.1	1.0934	+0.0000j	1.0934	0.0	
000					
MA.1	1.0578	+0.0000j	1.0578	0.0	
000					
MA.2	11.1424	+0.0000j	11.1424	0.0	
000					

In [28]:


```
#Forecasting the test data
```

In [25]:

```
fc, se, conf = fitted.forecast(321, alpha=0.05)
```

In [26]:

```
fc_series = pd.Series(fc, index=test_data.index)
lower_series = pd.Series(conf[:, 0], index=test_data.index)
upper_series = pd.Series(conf[:, 1], index=test_data.index)
# Plot
plt.figure(figsize=(10,5), dpi=100)
plt.plot(train_data, label='training data')
plt.plot(test_data, color = 'blue', label='Actual Stock Price')
plt.plot(fc_series, color = 'orange', label='Predicted Stock Price')
plt.fill_between(lower_series.index, lower_series, upper_series,
                 color='k', alpha=.10)
plt.title('ARCH CAPITAL GROUP Stock Price Prediction')
plt.xlabel('Time')
plt.ylabel('ARCH CAPITAL GROUP Stock Price')
plt.legend(loc='upper left', fontsize=8)
plt.show()
```



In [27]:

```
mse = mean_squared_error(test_data, fc)
print('MSE: '+str(mse))
mae = mean_absolute_error(test_data, fc)
print('MAE: '+str(mae))
rmse = math.sqrt(mean_squared_error(test_data, fc))
print('RMSE: '+str(rmse))
mape = np.mean(np.abs(fc - test_data)/np.abs(test_data))
print('MAPE: '+str(mape))
```

```
MSE: 0.015076603462322346
MAE: 0.11500989703236117
RMSE: 0.12278682120782484
MAPE: 0.02539744321032314
```

In [29]:

```
#Around 2.5% MAPE implies the model is about 97.5% accurate in predicting the next 15 observations.
```

In []: